



# CUDA **CUSPARSE Library**

---

PG-05329-032\_V01  
August, 2010

Published by  
NVIDIA Corporation  
2701 San Tomas Expressway  
Santa Clara, CA 95050

**Notice**

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS". NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication or otherwise under any patent or patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all information previously supplied. NVIDIA Corporation products are not authorized for use as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

**Trademarks**

NVIDIA, CUDA, and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the United States and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

**Copyright**

© 2005–2010 by NVIDIA Corporation. All rights reserved.

# Table of Contents

<b>1. CUSPARSE Library</b>	<b>5</b>
CUSPARSE Formats	6
Index Base Format	6
Sparse Vector Format	6
Matrix Formats	7
Dense Format	7
Coordinate Format (COO)	8
Compressed Sparse Row Format (CSR)	9
Compressed Sparse Column Format (CSC)	10
CUSPARSE Types	12
cusparseHandle_t	12
cusparseMatrixType_t	12
cusparseFillMode_t	13
cusparseDiagType_t	13
cusparseIndexBase_t	13
cusparseMatDescr_t	13
cusparseOperation_t	13
cusparseDirection_t	14
cusparseStatus_t	14
<b>2. CUSPARSE Functions</b>	<b>16</b>
CUSPARSE Helper Functions	17
cusparseCreate()	17
cusparseDestroy()	18
cusparseGetVersion()	18
cusparseSetKernelStream()	19
cusparseCreateMatDescr()	19
cusparseDestroyMatDescr()	20
cusparseSetMatType()	20
cusparseGetMatType()	20
cusparseSetMatFillMode()	21
cusparseGetMatFillMode()	21
cusparseSetMatDiagType()	21
cusparseGetMatDiagType()	22
cusparseSetMatIndexBase()	22
cusparseGetMatIndexBase()	23

Naming Convention for the Sparse Level Functions. . . . .	23
Sparse Level 1 Functions . . . . .	23
cusparse{S,D,C,Z}axpyi . . . . .	24
cusparse{S,D,C,Z}doti . . . . .	25
cusparse{C,Z}dotci . . . . .	27
cusparse{S,D,C,Z}gthr . . . . .	28
cusparse{S,D,C,Z}gthrz . . . . .	29
cusparse{S,D}roti . . . . .	31
cusparse{S,D,C,Z}sctr . . . . .	32
Sparse Level 2 Function . . . . .	33
cusparse{S,D,C,Z}csmv . . . . .	33
Sparse Level 3 Function . . . . .	36
cusparse{S,D,C,Z}csrmm . . . . .	36
Format Conversion Functions . . . . .	39
cusparse{S,D,C,Z}nnz . . . . .	39
cusparse{S,D,C,Z}dense2csr . . . . .	41
cusparse{S,D,C,Z}csr2dense . . . . .	43
cusparse{S,D,C,Z}dense2csc . . . . .	44
cusparse{S,D,C,Z}csc2dense . . . . .	46
cusparse{S,D,C,Z}csr2csc . . . . .	48
cusparseXcoo2csr . . . . .	50
cusparseXcsr2coo . . . . .	51
<b>A. CUSPARSE Library Example . . . . .</b>	<b>53</b>

## CHAPTER

# 1

## CUSPARSE Library

The NVIDIA® CUDA™ CUSPARSE library contains a set of basic linear algebra subroutines used for handling sparse matrices and is designed to be called from C or C<sup>++</sup>. These subroutines can be classified in four categories:

- ❑ Level 1 routines include operations between a vector in sparse format and a vector in dense format.
- ❑ Level 2 routines include operations between a matrix in sparse format and a vector in dense format.
- ❑ Level 3 routines include operations between a matrix in sparse format and a set of vectors (tall matrix) in dense format.
- ❑ Conversion routines that allow conversion between different matrix formats.

The library is written using the CUDA parallel programming model and takes advantage of the computational resources of the NVIDIA graphics processor (GPU). The CUSPARSE API assumes that the input and output data reside in GPU (device) memory, not in CPU (host) memory, unless CPU memory is specifically indicated by the string `HostPtr` being part of the parameter name of a function (for example, `*resultHostPtr` in `cusparse[S,D,C,Z]doti` on page 25).

It is the responsibility of the user to allocate memory and to copy data between GPU memory and CPU memory using standard CUDA runtime API routines, such as, `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`, and `cudaMemcpyAsync()`. (The CUDA runtime API is part of the CUDA Toolkit from NVIDIA.) The library is currently designed to run only on single-GPU systems; it does not auto-parallelize across multiple GPUs.

---

**Note:** The CUSPARSE library requires hardware with at least 1.1 compute capability. Please see *NVIDIA CUDA C Programming Guide*, Appendix A for the list of all compute capabilities.

---

## CUSPARSE Formats

CUSPARSE supports a [Sparse Vector Format](#) and [Matrix Formats](#).

### Index Base Format

The library supports zero- and one-based indexing.

### Sparse Vector Format

Sparse vectors are represented with two arrays.

- ❑ One data array contains all the non-zero values from the equivalent array in dense format.
- ❑ One integer index array contains the position of the corresponding non-zero value in the equivalent array in dense format.

For example, this  $7 \times 1$  dense vector can be stored as a one-based or a zero-based sparse vector.

$7 \times 1$  vector:  $\begin{bmatrix} 1.0 & 0.0 & 0.0 & 2.0 & 3.0 & 0.0 & 4.0 \end{bmatrix}$

One-based:  $\begin{bmatrix} 1.0 & 2.0 & 3.0 & 4.0 \\ 1 & 4 & 5 & 7 \end{bmatrix}$

Zero-based: 
$$\begin{bmatrix} 1.0 & 2.0 & 3.0 & 4.0 \\ 0 & 3 & 4 & 6 \end{bmatrix}$$

---

**Note:** It is assumed that the indices are provided in an increasing order and that each index appears only once.

---

## Matrix Formats

The matrix formats are the following:

- ❑ [Dense Format](#) on page 7
- ❑ [Coordinate Format \(COO\)](#) on page 8
- ❑ [Compressed Sparse Row Format \(CSR\)](#) on page 9
- ❑ [Compressed Sparse Column Format \(CSC\)](#) on page 10

### Dense Format

The dense matrix  $X$  is represented by the following parameters (assuming it is stored in column-major format in memory):

- ❑  $m$  (integer): the number of rows in the matrix.
- ❑  $n$  (integer): the number of columns in the matrix.
- ❑  $ldX$  (integer): the leading dimension of  $X$ , which must be greater than or equal to  $m$ . If  $ldX$  is greater than  $m$ , then  $X$  represents a sub-matrix of a larger  $ldX \times n$  matrix stored in memory.
- ❑  $X$  (pointer): points to the data array containing the matrix elements. It is assumed that enough storage is allocated for  $X$  to hold at least  $ldX * n$  matrix elements and that CUSPARSE library functions may access values outside of the  $m \times n$  sub-matrix, but will never overwrite them.

[Figure 1](#) on page 8 shows a schematic representation of the  $m \times n$  dense matrix  $X$  (shaded area) with leading dimension  $ldX$  greater than  $m$  and one-based indexing.




Figure 1.  $m \times n$  Dense Matrix  $X$  with  $ldX > m$

Please note that this format and notation is similar to the format and notation used in the NVIDIA CUDA CUBLAS library.

## Coordinate Format (COO)

The  $m \times n$  sparse matrix  $A$  is represented in COO format by the following parameters:

- ❑ `nnz` (integer): the number of non-zero elements in the matrix.
- ❑ `cooValA` (pointer): points to the data array of length `nnz` that holds all non-zero values of  $A$  in row-major format.
- ❑ `cooRowIndA` (pointer): points to the integer array of length `nnz` that contains the row indices of the corresponding elements in array `cooValA`.
- ❑ `cooColIndA` (pointer): points to the integer array of length `nnz` that contains the column indices of the corresponding elements in array `cooValA`.

---

**Note:** It is assumed that the indices are given in row-major format (first sorted by row indices and then within the same row by column indices) and that each pair of row and column indices appears only once.

---



Consider the following  $4 \times 5$  matrix A.

$$\begin{bmatrix} 1.0 & 4.0 & 0.0 & 0.0 & 0.0 \\ 0.0 & 2.0 & 3.0 & 0.0 & 0.0 \\ 5.0 & 0.0 & 0.0 & 7.0 & 8.0 \\ 0.0 & 0.0 & 9.0 & 0.0 & 6.0 \end{bmatrix}$$

This is how it is stored in COO zero-based format.

$$\begin{aligned} \text{cooValA} &= [1.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 5.0 \quad 7.0 \quad 8.0 \quad 9.0 \quad 6.0] \\ \text{cooRowIndA} &= [0 \quad 0 \quad 1 \quad 1 \quad 2 \quad 2 \quad 2 \quad 3 \quad 3] \\ \text{cooColIndA} &= [0 \quad 1 \quad 1 \quad 2 \quad 0 \quad 3 \quad 4 \quad 2 \quad 4] \end{aligned}$$

And this is the COO one-based format.

$$\begin{aligned} \text{cooValA} &= [1.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 5.0 \quad 7.0 \quad 8.0 \quad 9.0 \quad 6.0] \\ \text{cooRowIndA} &= [1 \quad 1 \quad 2 \quad 2 \quad 3 \quad 3 \quad 3 \quad 4 \quad 4] \\ \text{cooColIndA} &= [1 \quad 2 \quad 2 \quad 3 \quad 1 \quad 4 \quad 5 \quad 3 \quad 5] \end{aligned}$$

## Compressed Sparse Row Format (CSR)

The only difference between the COO and CSR formats is that the array containing the row indices is compressed in CSR format.

The  $m \times n$  sparse matrix A is represented in CSR format by the following parameters:

- ❑ `nnz` (integer): the number of non-zero elements in the matrix.
- ❑ `csrValA` (pointer): points to the data array of length `nnz` that holds all non-zero values of A in row-major format.
- ❑ `csrRowPtrA` (pointer): points to the integer array of length `m+1` that holds indices pointing to the array `csrColIndA/csrValA`. For the first `m` entries, `csrRowPtrA(i)` contains the index of the first non-zero element in the  $i^{\text{th}}$  row, while the last entry, `csrRowPtrA(m)`, contains `nnz+csrRowPtrA(0)`. In general, `csrRowPtrA(0)` is 0 or 1 depending on whether zero- or one-based format is used, respectively.

- ❑ `csrColIndA` (pointer): points to the integer array of length `nnz` that holds the column indices of the corresponding elements in `csrValA`.

---

**Note:** It is assumed that the indices are given in row-major format (first sorted by row indices and then within the same row by column indices) and that each pair of row and column indices appears only once.

---

Again, consider the  $4 \times 5$  matrix  $A$ .

$$\begin{bmatrix} 1.0 & 4.0 & 0.0 & 0.0 & 0.0 \\ 0.0 & 2.0 & 3.0 & 0.0 & 0.0 \\ 5.0 & 0.0 & 0.0 & 7.0 & 8.0 \\ 0.0 & 0.0 & 9.0 & 0.0 & 6.0 \end{bmatrix}$$

It is stored in CSR zero-based format as shown.

$$\begin{aligned} \text{csrValA} &= [1.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 5.0 \quad 7.0 \quad 8.0 \quad 9.0 \quad 6.0] \\ \text{csrRowPtrA} &= [0 \quad 2 \quad 4 \quad 7 \quad 9] \\ \text{csrColIndA} &= [0 \quad 1 \quad 1 \quad 2 \quad 0 \quad 3 \quad 4 \quad 2 \quad 4] \end{aligned}$$

This is the CSR one-based format.

$$\begin{aligned} \text{csrValA} &= [1.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 5.0 \quad 7.0 \quad 8.0 \quad 9.0 \quad 6.0] \\ \text{csrRowPtrA} &= [1 \quad 3 \quad 5 \quad 8 \quad 10] \\ \text{csrColIndA} &= [1 \quad 2 \quad 2 \quad 3 \quad 1 \quad 4 \quad 5 \quad 3 \quad 5] \end{aligned}$$

## Compressed Sparse Column Format (CSC)

The  $m \times n$  matrix  $A$  is represented in CSC format by the following parameters:

- ❑ `nnz` (integer): the number of non-zero elements in the matrix.
- ❑ `cscValA` (pointer): points to the data array of length `nnz` that holds all non-zero values of  $A$  in column-major format.
- ❑ `cscRowIndA` (pointer): points to the integer array of length `nnz` that holds the row indices of the corresponding elements in `cscValA`.

- `cscColPtrA` (pointer): points to the integer array of length  $n+1$  that holds indices pointing to array `cscRowIndA/cscValA`. For the first  $n$  entries, `cscColPtrA(i)` contains the index of the first non-zero element in the  $i^{\text{th}}$  column, while the last entry, `cscColPtrA(n)`, contains  $\text{nnz} + \text{cscColPtrA}(0)$ . In general, `cscColPtrA(0)` is 0 or 1 depending on whether zero- or one-based format is used, respectively.

---

**Note:** It is assumed that the indices are given in column-major format (first sorted by column indices and then within the same column by row indices) and that each pair of row and column indices appears only once.

---

Also note that matrix  $A$  in CSR format has exactly the same memory layout as its transpose in CSC format (and vice-versa).

Consider the  $4 \times 5$  matrix  $A$  one more time.

$$\begin{bmatrix} 1.0 & 4.0 & 0.0 & 0.0 & 0.0 \\ 0.0 & 2.0 & 3.0 & 0.0 & 0.0 \\ 5.0 & 0.0 & 0.0 & 7.0 & 8.0 \\ 0.0 & 0.0 & 9.0 & 0.0 & 6.0 \end{bmatrix}$$

The CSC zero-based storage format is below.

$$\begin{aligned} \text{cscValA} &= [1.0 \quad 5.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 9.0 \quad 7.0 \quad 8.0 \quad 6.0] \\ \text{cscRowIndA} &= [0 \quad 2 \quad 0 \quad 1 \quad 1 \quad 3 \quad 2 \quad 2 \quad 3] \\ \text{cscColPtrA} &= [0 \quad 2 \quad 4 \quad 6 \quad 7 \quad 9] \end{aligned}$$

And, this is the CSC one-based format.

$$\begin{aligned} \text{cscValA} &= [1.0 \quad 5.0 \quad 4.0 \quad 2.0 \quad 3.0 \quad 9.0 \quad 7.0 \quad 8.0 \quad 6.0] \\ \text{cscRowIndA} &= [1 \quad 3 \quad 1 \quad 2 \quad 2 \quad 4 \quad 3 \quad 3 \quad 4] \\ \text{cscColPtrA} &= [1 \quad 3 \quad 5 \quad 7 \quad 8 \quad 10] \end{aligned}$$

---

## CUSPARSE Types

The library supports the following data types: `float`, `double`, `cuComplex`, and `cuDoubleComplex`. The first two are standard C data types, the second two are exported from `cuComplex.h`.

The CUSPARSE library provides these types:

- ❑ [cusparseHandle\\_t](#) on page 12
- ❑ [cusparseMatrixType\\_t](#) on page 12
- ❑ [cusparseFillMode\\_t](#) on page 13
- ❑ [cusparseDiagType\\_t](#) on page 13
- ❑ [cusparseIndexBase\\_t](#) on page 13
- ❑ [cusparseMatDescr\\_t](#) on page 13
- ❑ [cusparseOperation\\_t](#) on page 13
- ❑ [cusparseDirection\\_t](#) on page 14
- ❑ [cusparseStatus\\_t](#) on page 14

### `cusparseHandle_t`

This is a pointer type to an opaque CUSPARSE context, which the user must initialize by calling `cusparseCreate()` prior to calling any other library function. The handle created and returned by `cusparseCreate()` must be passed to every CUSPARSE function.

### `cusparseMatrixType_t`

This type indicates the type of matrix stored in sparse storage.

```
typedef enum {  
    CUSPARSE_MATRIX_TYPE_GENERAL=0,  
    CUSPARSE_MATRIX_TYPE_SYMMETRIC=1,  
    CUSPARSE_MATRIX_TYPE_HERMITIAN=2,  
    CUSPARSE_MATRIX_TYPE_TRIANGULAR=3  
} cusparseMatrixType_t;
```

## cusparseFillMode\_t

This type indicates if the lower or upper part of a matrix is stored in sparse storage.

```
typedef enum {  
    CUSPARSE_FILL_MODE_LOWER=0,  
    CUSPARSE_FILL_MODE_UPPER=1  
} cusparseFillMode_t;
```

## cusparseDiagType\_t

This type indicates if the matrix diagonal entries are equal to one.

```
typedef enum {  
    CUSPARSE_DIAG_TYPE_NON_UNIT=0,  
    CUSPARSE_DIAG_TYPE_UNIT=1  
} cusparseDiagType_t;
```

## cusparseIndexBase\_t

This type indicates if the base of the matrix indices is zero or one.

```
typedef enum {  
    CUSPARSE_INDEX_BASE_ZERO=0,  
    CUSPARSE_INDEX_BASE_ONE=1  
} cusparseIndexBase_t;
```

## cusparseMatDescr\_t

This structure is used to describe the shape and properties of a matrix.

```
typedef struct {  
    cusparseMatrixType_t MatrixType;  
    cusparseFillMode_t FillMode;  
    cusparseDiagType_t DiagType;  
    cusparseIndexBase_t IndexBase;  
} cusparseMatDescr_t;
```

## cusparseOperation\_t

Indicates which operations need to be performed with the sparse matrix.

```
typedef enum {
    CUSPARSE_OPERATION_NON_TRANSPOSE=0,
    CUSPARSE_OPERATION_TRANSPOSE=1,
    CUSPARSE_OPERATION_CONJUGATE_TRANSPOSE=2
} cusparseOperation_t;
```

## cusparseDirection\_t

Indicates whether the elements of a matrix should be parsed by rows or by columns (regardless of row- or column-major storage format).

```
typedef enum {
    CUSPARSE_DIRECTION_ROW=0,
    CUSPARSE_DIRECTION_COLUMN=1
} cusparseDirection_t;
```

## cusparseStatus\_t

This is a status type returned by the library functions and can have the following defined values:

```
typedef enum{
    CUSPARSE_STATUS_SUCCESS=0,
    CUSPARSE_STATUS_NOT_INITIALIZED=1,
    CUSPARSE_STATUS_ALLOC_FAILED=2,
    CUSPARSE_STATUS_INVALID_VALUE=3,
    CUSPARSE_STATUS_ARCH_MISMATCH=4,
    CUSPARSE_STATUS_MAPPING_ERROR=5,
    CUSPARSE_STATUS_EXECUTION_FAILED=6,
    CUSPARSE_STATUS_INTERNAL_ERROR=7,
    CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED=8,
} cusparseStatus_t;
```

The status values are explained in the table below:

### CUSPARSE Status Definitions

#### CUSPARSE\_STATUS\_SUCCESS

The operation completed successfully.

#### CUSPARSE\_STATUS\_NOT\_INITIALIZED

The CUSPARSE library was not initialized. This is usually caused by the lack of a prior **cusparseCreate()** call.

To correct: call **cusparseCreate()** prior to the function call.

## CUSPARSE Status Definitions (continued)

### **CUSPARSE\_STATUS\_ALLOC\_FAILED**

Resource allocation failed inside the CUSPARSE library. This is usually caused by a **cudaMalloc()** failure.

To correct: prior to the function call, deallocate previously allocated memory as much as possible.

### **CUSPARSE\_STATUS\_INVALID\_VALUE**

An unsupported value or parameter was passed to the function (a negative vector size, for example).

To correct: ensure that all the parameters being passed have valid values.

### **CUSPARSE\_STATUS\_ARCH\_MISMATCH**

Function requires a feature absent from the device architecture; usually caused by the lack of support for atomic operations or double precision.

To correct: compile and run the application on a device with appropriate compute capability, which is 1.1 for 32-bit atomic operations and 1.3 for double precision.

### **CUSPARSE\_STATUS\_MAPPING\_ERROR**

An access to GPU memory space failed, which is usually caused by a failure to bind a texture.

To correct: prior to the function call, unbind any previously bound textures.

### **CUSPARSE\_STATUS\_EXECUTION\_FAILED**

The GPU program failed to execute. This is often caused by a launch failure of the kernel on the GPU, which can be caused by multiple reasons.

To correct: check that the hardware, an appropriate version of the driver, and the CUSPARSE library are correctly installed.

### **CUSPARSE\_STATUS\_INTERNAL\_ERROR**

An internal CUSPARSE operation failed. This error is usually caused by a **cudaMemcpyAsync()** failure.

To correct: check that the hardware, an appropriate version of the driver, and the CUSPARSE library are correctly installed.

### **CUSPARSE\_STATUS\_MATRIX\_TYPE\_NOT\_SUPPORTED**

The matrix type is not supported by this function. This is usually caused by passing an invalid matrix descriptor to the function.

To correct: check that the fields in **cusparseMatDescr\_t \*descrA** were set correctly.

## CUSPARSE Functions

This chapter discusses the CUSPARSE functions, which are divided into five groups, and the naming convention used for Sparse Level 1, Level 2, and Level 3 functions.

- ❑ [CUSPARSE Helper Functions](#) on page 17
- ❑ [Naming Convention for the Sparse Level Functions](#) on page 23
- ❑ [Sparse Level 1 Functions](#) on page 23
- ❑ [Sparse Level 2 Function](#) on page 33
- ❑ [Sparse Level 3 Function](#) on page 36
- ❑ [Format Conversion Functions](#) on page 39



## CUSPARSE Helper Functions

The CUSPARSE helper functions are as follows:

- ❑ [cusparsesetCreate\(\)](#) on page 17
- ❑ [cusparsesetDestroy\(\)](#) on page 18
- ❑ [cusparsesetGetVersion\(\)](#) on page 18
- ❑ [cusparsesetSetKernelStream\(\)](#) on page 19
- ❑ [cusparsesetCreateMatDescr\(\)](#) on page 19
- ❑ [cusparsesetDestroyMatDescr\(\)](#) on page 20
- ❑ [cusparsesetSetMatType\(\)](#) on page 20
- ❑ [cusparsesetGetMatType\(\)](#) on page 20
- ❑ [cusparsesetSetMatFillMode\(\)](#) on page 21
- ❑ [cusparsesetGetMatFillMode\(\)](#) on page 21
- ❑ [cusparsesetSetMatDiagType\(\)](#) on page 21
- ❑ [cusparsesetGetMatDiagType\(\)](#) on page 22
- ❑ [cusparsesetSetMatIndexBase\(\)](#) on page 22
- ❑ [cusparsesetGetMatIndexBase\(\)](#) on page 23

### cusparsesetCreate()

```
cusparsesetStatus_t
cusparsesetCreate( cusparsesetHandle_t *handle )
```

Initializes the CUSPARSE library and creates a handle on the CUSPARSE context. This function must be called before any other CUSPARSE API function is invoked. It allocates hardware resources necessary for accessing the GPU.

---

**Note:** The CUSPARSE library requires hardware with at least 1.1 compute capability. Please see *NVIDIA CUDA C Programming Guide*, Appendix A for the list of all compute capabilities.

---

#### Output

---

<b>handle</b>	initialized pointer to a CUSPARSE context
---------------	---

---

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	CUSPARSE library initialized successfully
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	if an error with CUDA or the hardware setup has been detected
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if device compute capability is less than 1.1

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseDestroy()

```
cusparseStatus_t  
cusparseDestroy( cusparseHandle_t handle )
```

Releases CPU side resources used by the CUSPARSE library. The release of GPU side resources may be deferred until the application shuts down.

## Input

<b>handle</b>	handle to a CUSPARSE context
---------------	------------------------------

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	CUSPARSE library shut down successfully
--------------------------------	---

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseGetVersion()

```
cusparseStatus_t  
cusparseGetVersion(  
    cusparseHandle_t handle, int *version )
```

Returns the version number of the CUSPARSE library.

## Input

<b>handle</b>	handle to a CUSPARSE context
---------------	------------------------------

## Output

<b>version</b>	integer version number of the library
----------------	---------------------------------------

Status Returned<sup>i</sup>

**CUSPARSE\_STATUS\_SUCCESS**

CUSPARSE library version was returned successfully

**CUSPARSE\_STATUS\_NOT\_INITIALIZED**

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseSetKernelStream()

**cusparseStatus\_t**

**cusparseSetKernelStream(**

**cusparseHandle\_t handle, cudaStream\_t streamId )**

Sets the CUSPARSE stream in which the kernels will run.

Input

**handle** handle to a CUSPARSE context

Status Returned<sup>i</sup>

**CUSPARSE\_STATUS\_SUCCESS**

if stream provided has been set properly

**CUSPARSE\_STATUS\_NOT\_INITIALIZED**

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseCreateMatDescr()

**cusparseStatus\_t**

**cusparseCreateMatDescr( cusparseMatDescr\_t \*descrA )**

Initializes the **MatrixType** and **IndexBase** fields of the matrix descriptor to the default values **CUSPARSE\_MATRIX\_TYPE\_GENERAL** and **CUSPARSE\_INDEX\_BASE\_ZERO**, while leaving other fields uninitialized.

Input

**descrA** descriptor of the matrix A

Status Returned

**CUSPARSE\_STATUS\_SUCCESS**

matrix descriptor initialized successfully

**CUSPARSE\_STATUS\_ALLOC\_FAILED**

if resources could not be allocated for the matrix descriptor

## cusparseDestroyMatDescr()

**cusparseStatus\_t**

**cusparseDestroyMatDescr( cusparseMatDescr\_t descrA )**

Releases the memory allocated for the matrix descriptor.

Input

descrA	descriptor of the matrix A
--------	----------------------------

Status Returned

<b>CUSPARSE_STATUS_SUCCESS</b>	matrix descriptor destroyed successfully
--------------------------------	--

## cusparseSetMatType()

**cusparseStatus\_t**

**cusparseSetMatType( cusparseMatDescr\_t descrA, cusparseMatrixType\_t type )**

Sets the **MatrixType** field of the matrix descriptor descrA.

Input

type	one of the enumerated matrix types
------	------------------------------------

Output

descrA	descriptor of the matrix A
--------	----------------------------

Status Returned

<b>CUSPARSE_STATUS_SUCCESS</b>	<b>MatrixType</b> field was set successfully
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if invalid value was passed in the type parameter

## cusparseGetMatType()

**cusparseMatrixType\_t**

**cusparseGetMatType( const cusparseMatDescr\_t descrA )**

Returns the **MatrixType** field of the matrix descriptor descrA.

Input

descrA	descriptor of the matrix A
--------	----------------------------

Status Returned

one of the enumerated matrix types
------------------------------------

## cusparseSetMatFillMode()

```
cusparseStatus_t
cusparseSetMatFillMode(
    cusparseMatDescr_t descrA, cusparseFillMode_t fillMode )
```

Sets the **FillMode** field of the matrix descriptor descrA.

Input

---

fillMode	one of the enumerated matrix types
----------	------------------------------------

---

Output

---

descrA	descriptor of the matrix A
--------	----------------------------

---

Status Returned

---

CUSPARSE_STATUS_SUCCESS	<b>FillMode</b> field was set successfully
CUSPARSE_STATUS_INVALID_VALUE	if invalid value was passed in the fillMode parameter

---

## cusparseGetMatFillMode()

```
cusparseFillMode_t
cusparseGetMatFillMode( const cusparseMatDescr_t descrA )
```

Returns the **FillMode** field of the matrix descriptor descrA.

Input

---

descrA	descriptor of the matrix A
--------	----------------------------

---

Status Returned

---

one of the enumerated fill-in types
-------------------------------------

---

## cusparseSetMatDiagType()

```
cusparseStatus_t
cusparseSetMatDiagType(
    cusparseMatDescr_t descrA, cusparseDiagType_t diagType )
```

Sets the **DiagType** field of the matrix descriptor descrA.

Input

---

diagType	one of the enumerated diagonal modes
----------	--------------------------------------

---

## Output

---

`descrA`   descriptor of the matrix A

---

## Status Returned

<b>CUSPARSE_STATUS_SUCCESS</b>	<b>DiagType</b> field was set successfully
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if invalid value was passed in the <b>diagType</b> parameter

---

## cusparseGetMatDiagType()

**cusparseDiagType\_t**

**cusparseGetMatDiagType( const cusparseMatDescr\_t descrA )**

Returns the **DiagType** field of the matrix descriptor `descrA`.

## Input

---

`descrA`   descriptor of the matrix A

---

## Status Returned

---

one of the enumerated diagonal modes

---

## cusparseSetMatIndexBase()

**cusparseStatus\_t**

**cusparseSetMatIndexBase(**

**cusparseMatDescr\_t descrA, cusparseIndexBase\_t base )**

Sets the **IndexBase** field of the matrix descriptor `descrA`.

## Input

---

`base`      one of the enumerated index base modes

---

## Output

---

`descrA`   descriptor of the matrix A

---

## Status Returned

<b>CUSPARSE_STATUS_SUCCESS</b>	<b>IndexBase</b> field was set successfully
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if invalid value was passed in the <b>base</b> parameter

---

## cusparseGetMatIndexBase()

**cusparseIndexBase\_t**

**cusparseGetMatIndexBase( const cusparseMatDescr\_t descrA )**

Returns the **IndexBase** field of the matrix descriptor **descrA**.

Input

---

**descrA**    descriptor of the matrix A

---

Status Returned

---

one of the enumerated index base modes

---

## Naming Convention for the Sparse Level Functions

Most of the CUSPARSE functions are available for data types **float**, **double**, **cuComplex**, and **cuDoubleComplex**. The Sparse Level 1, Level 2, and Level 3 functions follow this naming convention:

**cusparse**<T>[<sparse data format>]<operation>[<sparse data format>]

with <T> being **S**, **D**, **C**, **Z**, or **X** corresponding to the data types **float**, **double**, **cuComplex**, **cuDoubleComplex**, or no type, respectively. All these functions have the return type **cusparseStatus\_t**.

## Sparse Level 1 Functions

The Level 1 functions are the following:

- ❑ [cusparse{S,D,C,Z}axpyi](#) on page 24
- ❑ [cusparse{S,D,C,Z}doti](#) on page 25
- ❑ [cusparse{C,Z}dotci](#) on page 27
- ❑ [cusparse{S,D,C,Z}gthr](#) on page 28
- ❑ [cusparse{S,D,C,Z}gthrz](#) on page 29
- ❑ [cusparse{S,D}roti](#) on page 31
- ❑ [cusparse{S,D,C,Z}sctr](#) on page 32

## cusparse{S,D,C,Z}axpyi

```

cusparseStatus_t
cusparseSaxpyi(
    cusparseHandle_t handle, int nnz,
    float alpha, const float *xVal,
    const int *xInd, float *y,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseDaxpyi(
    cusparseHandle_t handle, int nnz,
    double alpha, const double *xVal,
    const int *xInd, double *y,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseCaxpyi(
    cusparseHandle_t handle, int nnz,
    cuComplex alpha, const cuComplex *xVal,
    const int *xInd, cuComplex *y,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseZaxpyi(
    cusparseHandle_t handle, int nnz,
    cuDoubleComplex alpha, const cuDoubleComplex *xVal,
    const int *xInd, cuDoubleComplex *y,
    cusparseIndexBase_t idxBase )

```

Multiplies the vector  $x$  in sparse format by the constant  $\alpha$  and adds the result to the vector  $y$  in dense format; that is, it overwrites  $y$  with  $\alpha * x + y$ .

For  $i = 0$  to  $nnz-1$

$$y[xInd[i] - idxBase] = y[xInd[i] - idxBase] + \alpha * xVal[i]$$

Input

---

handle	handle to a CUSPARSE context
nnz	number of elements of the vector $x$
alpha	constant multiplier
xVal	nnz non-zero values of vector $x$
xInd	nnz indices corresponding to non-zero values of vector $x$



Input (continued)

y	initial vector in dense format
idxBase	CUSPARSE_INDEX_BASE_ZERO or CUSPARSE_INDEX_BASE_ONE

Output

y	result (unchanged if nnz == 0)
---	--------------------------------

Status Returned<sup>i</sup>

CUSPARSE_STATUS_SUCCESS	
CUSPARSE_STATUS_NOT_INITIALIZED	
CUSPARSE_STATUS_INVALID_VALUE	if idxBase is neither CUSPARSE_INDEX_BASE_ZERO nor CUSPARSE_INDEX_BASE_ONE
CUSPARSE_STATUS_ARCH_MISMATCH	if the D or Z variants of the function were invoked on a device that does not support double precision.
CUSPARSE_STATUS_EXECUTION_FAILED	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{ S,D,C,Z }doti

```
cusparseStatus_t
cusparseSdoti(
    cusparseHandle_t handle, int nnz,
    const float *xVal, const int *xInd,
    const float *y, float *resultHostPtr,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseDdoti(
    cusparseHandle_t handle, int nnz,
    const double *xVal, const int *xInd,
    const float *y, double *resultHostPtr,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseCdoti(
    cusparseHandle_t handle, int nnz,
    const cuComplex *xVal, const int *xInd,
    const float *y, cuComplex *resultHostPtr,
    cusparseIndexBase_t idxBase )
```

```

cusparseStatus_t
cusparseZdoti(
    cusparseHandle_t handle, int nnz,
    const cuDoubleComplex *xVal, const int *xInd,
    const float *y, cuDoubleComplex *resultHostPtr,
    cusparseIndexBase_t idxBase )

```

Returns the dot product of a vector x in sparse format and vector y in dense format.

For  $i = 0$  to  $nnz-1$

```
resultHostPtr += xVal[i] * y[xInd[i - idxBase]]
```

#### Input

handle	handle to a CUSPARSE context
nnz	number of elements of the vector x
xVal	nnz non-zero values of vector x
xInd	nnz indices corresponding to non-zero values of vector x
y	vector in dense format
resultHostPtr	<b>pointer where to write the result in host memory</b>
idxBase	<b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

#### Output

resultHostPtr	<b>updated host memory with dot product</b> <b>(zero if <math>nnz == 0</math>)</b>
---------------	---

#### Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if idxBase is neither <b>CUSPARSE_INDEX_BASE_ZERO</b> nor <b>CUSPARSE_INDEX_BASE_ONE</b>
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the D or Z variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_INTERNAL_ERROR</b>	

<sup>i</sup>. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{C,Z}dotci

```

cusparseStatus_t
cusparseCdotci(
    cusparseHandle_t handle, int nnz,
    const cuComplex *xVal, const int *xInd,
    const cuComplex *y, cuComplex *resultHostPtr,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseZdotci(
    cusparseHandle_t handle, int nnz,
    const cuDoubleComplex *xVal, const int *xInd,
    const cuComplex *y, cuDoubleComplex *resultHostPtr,
    cusparseIndexBase_t idxBase )

```

Returns the dot product of a complex conjugate of vector x in sparse format and complex vector y in dense format.

It computes the sum for  $i = 0$  to  $nnz-1$  of

$$\text{resultHostPtr} += \overline{xVal[i]} * y[xInd[i - idxBase]]$$

### Input

handle	handle to a CUSPARSE context
nnz	number of elements of the vector x
xVal	nnz non-zero values of vector x
xInd	nnz indices corresponding to non-zero values of vector x
y	vector in dense format
resultHostPtr	pointer where to write the result in host memory
idxBase	CUSPARSE_INDEX_BASE_ZERO or CUSPARSE_INDEX_BASE_ONE

### Output

resultHostPtr	updated host memory with dot product (zero if $nnz == 0$ )
---------------	---

### Status Returned<sup>i</sup>

CUSPARSE_STATUS_SUCCESS	
CUSPARSE_STATUS_NOT_INITIALIZED	
CUSPARSE_STATUS_ALLOC_FAILED	
CUSPARSE_STATUS_INVALID_VALUE	if idxBase is neither CUSPARSE_INDEX_BASE_ZERO nor CUSPARSE_INDEX_BASE_ONE

Status Returned<sup>i</sup> (continued)

<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the D or Z variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_INTERNAL_ERROR</b>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}gthr

```

cusparseStatus_t
cusparseSgthr(
    cusparseHandle_t handle, int nnz,
    const float *y, float *xVal,
    const int *xInd, cusparseIndexBase_t idxBase )
cusparseStatus_t
cusparseDgthr(
    cusparseHandle_t handle, int nnz,
    const double *y, double *xVal,
    const int *xInd, cusparseIndexBase_t idxBase )
cusparseStatus_t
cusparseCgthr(
    cusparseHandle_t handle, int nnz,
    const cuComplex *y, cuComplex *xVal,
    const int *xInd, cusparseIndexBase_t idxBase )
cusparseStatus_t
cusparseZgthr(
    cusparseHandle_t handle, int nnz,
    const cuDoubleComplex *y, cuDoubleComplex *xVal,
    const int *xInd, cusparseIndexBase_t idxBase )

```

Gathers the elements of the vector *y* listed by the index array *xInd* into the array *xVal*.

### Input

<i>handle</i>	handle to a CUSPARSE context
<i>nnz</i>	number of elements of the vector <i>x</i>
<i>y</i>	vector in dense format, of size greater than or equal to $\max(xInd) - idxBase + 1$

## Input (continued)

<code>xVal</code>	pre-allocated array in device memory of size greater than or equal to <code>nnz</code>
<code>xInd</code>	<code>nnz</code> indices corresponding to non-zero values of vector <code>x</code>
<code>idxBase</code>	<code>CUSPARSE_INDEX_BASE_ZERO</code> or <code>CUSPARSE_INDEX_BASE_ONE</code>

## Output

<code>xVal</code>	updated vector of <code>nnz</code> elements (unchanged if <code>nnz == 0</code> )
-------------------	---

Status Returned<sup>i</sup>

<code>CUSPARSE_STATUS_SUCCESS</code>	
<code>CUSPARSE_STATUS_NOT_INITIALIZED</code>	
<code>CUSPARSE_STATUS_INVALID_VALUE</code>	if <code>idxBase</code> is neither <code>CUSPARSE_INDEX_BASE_ZERO</code> nor <code>CUSPARSE_INDEX_BASE_ONE</code>
<code>CUSPARSE_STATUS_ARCH_MISMATCH</code>	if the <code>D</code> or <code>Z</code> variants of the function were invoked on a device that does not support double precision.
<code>CUSPARSE_STATUS_EXECUTION_FAILED</code>	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}gthrz

```

cusparseStatus_t
cusparseSgthrz(
    cusparseHandle_t handle, int nnz, float *y,
    float *xVal, const int *xInd,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseDgthrz(
    cusparseHandle_t handle, int nnz, double *y,
    double *xVal, const int *xInd,
    cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseCgthrz(
    cusparseHandle_t handle, int nnz, cuComplex *y,
    cuComplex *xVal, const int *xInd,
    cusparseIndexBase_t idxBase )

```

```

cusparseStatus_t
cusparseZgthrz(
    cusparseHandle_t handle, int nnz, cuDoubleComplex *y,
    cuDoubleComplex *xVal, const int *xInd,
    cusparseIndexBase_t idxBase )

```

Gathers the elements of the vector *y* listed by the index array *xInd* into the vector *x*, and zeroes those elements in the vector *y*.

#### Input

<i>handle</i>	handle to a CUSPARSE context
<i>nnz</i>	number of elements of the vector <i>x</i>
<i>y</i>	vector in dense format, of size greater than or equal to $\max(xInd) - idxBase + 1$
<i>xVal</i>	nnz non-zero values of vector <i>x</i>
<i>xInd</i>	nnz indices corresponding to non-zero values of vector <i>x</i>
<i>idxBase</i>	<b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

#### Output

<i>xVal</i>	nnz non-zero values of vector <i>x</i> (unchanged if $nnz == 0$ )
<i>xInd</i>	nnz indices corresponding to non-zero values of vector <i>x</i> (unchanged if $nnz == 0$ )
<i>y</i>	vector in dense format (unchanged if $nnz == 0$ )

#### Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if <i>idxBase</i> is neither <b>CUSPARSE_INDEX_BASE_ZERO</b> nor <b>CUSPARSE_INDEX_BASE_ONE</b>
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D}roti

```

cusparseStatus_t
cusparseSroti(
    cusparseHandle_t handle, int nnz, float *xVal,
    const int *xInd, float *y, float c,
    float s, cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseDroti(
    cusparseHandle_t handle, int nnz, double *xVal,
    const int *xInd, double *y, double c,
    double s, cusparseIndexBase_t idxBase )

```

Applies Givens rotation, defined by values  $c$  and  $s$ , to vectors  $x$  in sparse and  $y$  in dense format.

For  $i = 0$  to  $nnz-1$

```

    y[xInd[i] - idxBase] = c * y[xInd[i] - idxBase] - s * xVal[i]
    x[i] = c * xVal[i] + s * y[xInd[i] - idxBase]

```

### Input

handle	handle to a CUSPARSE context
nnz	number of elements of the vector $x$
xVal	nnz non-zero values of vector $x$
xInd	nnz indices corresponding to non-zero values of vector $x$
y	vector in dense format
c	scalar
s	scalar
idxBase	<b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

### Output

xVal	updated nnz non-zero values of vector $x$ (unchanged if $nnz == 0$ )
xInd	updated nnz indices corresponding to non-zero values of vector $x$ (unchanged if $nnz == 0$ )
y	updated vector in dense format (unchanged if $nnz == 0$ )

### Status Returned<sup>i</sup>

```

CUSPARSE_STATUS_SUCCESS
CUSPARSE_STATUS_NOT_INITIALIZED

```

Status Returned<sup>i</sup> (continued)

<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if idxBase is neither <b>CUSPARSE_INDEX_BASE_ZERO</b> nor <b>CUSPARSE_INDEX_BASE_ONE</b>
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}sctr

```

cusparseStatus_t
cusparseSsctr(
    cusparseHandle_t handle, int nnz,
    const float *xVal, const int *xInd,
    float *y, cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseDsctr(
    cusparseHandle_t handle, int nnz,
    const double *xVal, const int *xInd,
    double *y, cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseCsctr(
    cusparseHandle_t handle, int nnz,
    const cuComplex *xVal, const int *xInd,
    cuComplex *y, cusparseIndexBase_t idxBase )

cusparseStatus_t
cusparseZsctr(
    cusparseHandle_t handle, int nnz,
    const cuDoubleComplex *xVal, const int *xInd,
    cuDoubleComplex *y, cusparseIndexBase_t idxBase )

```

Scatters the vector *x* in sparse format into the vector *y* in dense format. It modifies only the elements of *y* whose indices are listed in the array *xInd*.

Input

handle	handle to a CUSPARSE context
nnz	number of elements of the vector <i>x</i>



## Input (continued)

xVal	nnz non-zero values of vector x
xInd	nnz indices corresponding to non-zero values of vector x
y	pre-allocated vector in dense format, of size greater than or equal to $\max(\text{xInd}) - \text{idxBase} + 1$
idxBase	<b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

## Output

y	updated vector in dense format (unchanged if $\text{nnz} == 0$ )
---	--

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if idxBase is neither <b>CUSPARSE_INDEX_BASE_ZERO</b> nor <b>CUSPARSE_INDEX_BASE_ONE</b>
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the D or Z variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

## Sparse Level 2 Function

There is one Level 2 function.

### cusparse{S,D,C,Z}csrmmv

```
cusparseStatus_t
cusparseScsrmmv(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, float alpha,
    const cusparseMatDescr_t *descrA,
    const float *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const float *x, float beta,
    float *y )
```

```

cusparseStatus_t
cusparseDcsrmmv(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, double alpha,
    const cusparseMatDescr_t *descrA,
    const double *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const double *x, double beta,
    double *y )

cusparseStatus_t
cusparseCcsrmmv(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, cuComplex alpha,
    const cusparseMatDescr_t *descrA,
    const cuComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const cuComplex *x, cuComplex beta,
    cuComplex *y )

cusparseStatus_t
cusparseZcsrmmv(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, cuDoubleComplex alpha,
    const cusparseMatDescr_t *descrA,
    const cuDoubleComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const cuDoubleComplex *x, cuDoubleComplex beta,
    cuDoubleComplex *y )

```

Performs one of the matrix-vector operations

$$y = \alpha * \text{op}(A) * x + \beta * y,$$

where  $\text{op}(A) = A$ ,  $\text{op}(A) = A^T$ , or  $\text{op}(A) = A^H$ .

A is an  $m \times n$  matrix in CSR format, defined by the three arrays `csrValA`, `csrRowPtrA`, and `csrColIndA`; `alpha` and `beta` are scalars, and `x` and `y` are vectors in dense format.

Input

<code>handle</code>	handle to a CUSPARSE context
<code>transA</code>	specifies $\text{op}(A)$ . If <code>transA == CUSPARSE_OPERATION_NON_TRANSPOSE</code> , $\text{op}(A) = A$ . This is the only operation supported at the moment.

## Input (continued)

<code>m</code>	specifies the number of rows of matrix A; <code>m</code> must be at least zero
<code>n</code>	specifies the number of columns of matrix A; <code>n</code> must be at least zero
<code>alpha</code>	scalar multiplier applied to $\text{op}(A) * x$
<code>descrA</code>	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
<code>csrValA</code>	array of <code>nnz</code> elements, where <code>nnz</code> is the number of non-zero elements and can be obtained from $\text{csrRowPtrA}(\text{m}) - \text{csrRowPtrA}(0)$
<code>csrRowPtrA</code>	array of <code>m+1</code> index elements
<code>csrColIndA</code>	array of <code>nnz</code> column indices
<code>x</code>	vector of <code>n</code> elements if $\text{op}(A) = A$ , and <code>m</code> elements if $\text{op}(A) = A^T$ or $\text{op}(A) = A^H$
<code>beta</code>	scalar multiplier applied to <code>y</code> . If <code>beta</code> is zero, <code>y</code> does not have to be a valid input.
<code>y</code>	vector of <code>m</code> elements if $\text{op}(A) = A$ , and <code>n</code> elements if $\text{op}(A) = A^T$ or $\text{op}(A) = A^H$

## Output

<code>y</code>	updated according to $y = \text{alpha} * \text{op}(A) * x + \text{beta} * y$
----------------	--

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_INTERNAL_ERROR</b>	
<b>CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED</b>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

---

## Sparse Level 3 Function

There is one Level 3 function.

### `cusparse{S,D,C,Z}csrmm`

```
cusparseStatus_t
cusparseScsrmm(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, int k, float alpha,
    const cusparseMatDescr_t *descrA,
    const float *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const float *B, int ldb,
    float beta, float *C, int ldc )

cusparseStatus_t
cusparseDcsrmm(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, int k, double alpha,
    const cusparseMatDescr_t *descrA,
    const double *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const double *B, int ldb,
    double beta, double *C, int ldc )

cusparseStatus_t
cusparseCcsrmm(
    cusparseHandle_t handle, cusparseOperation_t transA,
    int m, int n, int k, cuComplex alpha,
    const cusparseMatDescr_t *descrA,
    const cuComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const cuComplex *B, int ldb,
    cuComplex beta, cuComplex *C, int ldc )
```

```

cusparsesStatus_t
cusparsesZcsrmm(
    cusparsesHandle_t handle, cusparsesOperation_t transA,
    int m, int n, int k, cuDoubleComplex alpha,
    const cusparsesMatDescr_t *descrA,
    const cuDoubleComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    const cuDoubleComplex *B, int ldb,
    cuDoubleComplex beta, cuDoubleComplex *C, int ldc )

```

Performs one of these matrix-matrix operations:

$$C = \alpha * \text{op}(A) * B + \beta * C,$$

where  $\text{op}(A) = A$ ,  $\text{op}(A) = A^T$ , or  $\text{op}(A) = A^H$ ;

and  $\alpha$  and  $\beta$  are scalars.  $B$  and  $C$  are dense matrices stored in column-major format,  $A$  is an  $m \times k$  matrix in CSR format, defined by the three arrays  $\text{csrValA}$ ,  $\text{csrRowPtrA}$  and  $\text{csrColIndA}$ .

#### Input

handle	handle to a CUSPARSE context
transA	specifies $\text{op}(A)$ . If $\text{transA} == \text{CUSPARSE\_OPERATION\_NON\_TRANPOSE}$ , $\text{op}(A) = A$ . This is the only operation supported at the moment.
m	number of rows of matrix $A$ ; $m$ must be at least zero.
n	number of columns of matrices $B$ and $C$ ; $n$ must be at least zero.
k	number of columns of matrix $A$ ; $k$ must be at least zero.
alpha	scalar multiplier applied to $\text{op}(A) * B$
descrA	descriptor of matrix $A$ . The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
csrValA	array of $\text{nnz}$ elements, where $\text{nnz}$ is the number of non-zero elements and can be obtained from $\text{csrRowPtrA}(m) - \text{csrRowPtrA}(0)$
csrRowPtrA	array of $m+1$ index elements
csrColIndA	array of $\text{nnz}$ column indices
B	array of dimension $(\text{ldb}, n)$

## Input (continued)

ldb	leading dimension of B. It must be at least $\max(1, k)$ if $\text{op}(A) = A$ , and at least $\max(1, m)$ if $\text{op}(A) = A^T$ or $\text{op}(A) = A^H$ .
beta	scalar multiplier applied to C. If beta is zero, C does not have to be a valid input.
C	array of dimension (ldc, n)
ldc	leading dimension of C. It must be at least $\max(1, m)$ if $\text{op}(A) = A$ , and at least $\max(1, k)$ if $\text{op}(A) = A^T$ or $\text{op}(A) = A^H$ .

## Output

C	updated according to $\alpha * \text{op}(A) * B + \beta * C$
---	--

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_INTERNAL_ERROR</b>	
<b>CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED</b>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

## Format Conversion Functions

The format conversion functions are listed below:

- ❑ [cusparse{S,D,C,Z}nnz](#) on page 39
- ❑ [cusparse{S,D,C,Z}dense2csr](#) on page 41
- ❑ [cusparse{S,D,C,Z}csr2dense](#) on page 43
- ❑ [cusparse{S,D,C,Z}dense2csc](#) on page 44
- ❑ [cusparse{S,D,C,Z}csc2dense](#) on page 46
- ❑ [cusparse{S,D,C,Z}csr2csc](#) on page 48
- ❑ [cusparseXcoo2csr](#) on page 50
- ❑ [cusparseXcsr2coo](#) on page 51

### `cusparse{S,D,C,Z}nnz`

```

cusparseStatus_t
cusparseSnnz(
    cusparseHandle_t handle, cusparseDirection_t dirA,
    int m, int n, const cusparseMatDescr_t *descrA,
    const float *A, int lda, int *nnzPerVector,
    int *nnzHostPtr )

cusparseStatus_t
cusparseDnnz(
    cusparseHandle_t handle, cusparseDirection_t dirA,
    int m, int n, const cusparseMatDescr_t *descrA,
    const double *A, int lda, int *nnzPerVector,
    int *nnzHostPtr )

cusparseStatus_t
cusparseCnnz(
    cusparseHandle_t handle, cusparseDirection_t dirA,
    int m, int n, const cusparseMatDescr_t *descrA,
    const cuComplex *A, int lda, int *nnzPerVector,
    int *nnzHostPtr )

```

```

cusparseStatus_t
cusparseZnnz(
    cusparseHandle_t handle, cusparseDirection_t dirA,
    int m, int n, const cusparseMatDescr_t *descrA,
    const cuDoubleComplex *A, int lda, int *nnzPerVector,
    int *nnzHostPtr )

```

Computes the number of non-zero elements per row or column and the total number of non-zero elements.

#### Input

handle	handle to a CUSPARSE context
dirA	<b>CUSPARSE_DIRECTION_ROW</b> or <b>CUSPARSE_DIRECTION_COLUMN</b> indicates whether to count the number of non-zero elements per row or per column, respectively.
m	number of rows of the matrix A; m must be at least zero.
n	number of columns of matrix A; n must be at least zero.
descrA	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
A	array of dimension (lda, n)
lda	leading dimension of A
nnzPerVector	array of integers of size $m \times n$ to be filled
nnzHostPtr	pointer (in the host memory) to an integer to be filled

#### Output

nnzPerVector	array of size m or n, containing number of non-zero elements per row or column, respectively
nnzHostPtr	pointer (in the host memory) to an integer containing the total number of non-zero elements

#### Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU



Status Returned<sup>i</sup> (continued)

---

CUSPARSE\_STATUS\_INTERNAL\_ERROR

---

CUSPARSE\_STATUS\_MATRIX\_TYPE\_NOT\_SUPPORTED

---

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}dense2csr

```
cusparseStatus_t
cusparseSdense2csr(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const float *A, int lda,
    const int *nnzPerRow, float *csrValA,
    int *csrRowPtrA, int *csrColIndA )

cusparseStatus_t
cusparseDdense2csr(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const double *A, int lda,
    const int *nnzPerRow, double *csrValA,
    int *csrRowPtrA, int *csrColIndA )

cusparseStatus_t
cusparseCdense2csr(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const cuComplex *A, int lda,
    const int *nnzPerRow, cuComplex *csrValA,
    int *csrRowPtrA, int *csrColIndA )

cusparseStatus_t
cusparseZdense2csr(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const cuDoubleComplex *A, int lda,
    const int *nnzPerRow, cuDoubleComplex *csrValA,
    int *csrRowPtrA, int *csrColIndA )
```

Converts the matrix A in dense format into a matrix in CSR format. All the parameters are pre-allocated by the user, and the arrays are filled

in based on `nnzPerRow` (which can be pre-computed with `cusparseset{S,D,C,Z}nnz()`).

#### Input

<code>handle</code>	handle to a CUSPARSE context
<code>m</code>	number of rows of the matrix A; <code>m</code> must be at least zero.
<code>n</code>	number of columns of matrix A; <code>n</code> must be at least zero.
<code>descrA</code>	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
<code>A</code>	array of dimension <code>(lda, n)</code>
<code>lda</code>	leading dimension of A
<code>nnzPerRow</code>	array of size <code>m</code> containing the number of non-zero elements per row
<code>csrValA</code>	array of <code>nnz</code> elements to be filled
<code>csrRowPtrA</code>	array of <code>m+1</code> index elements
<code>csrColIndA</code>	array of <code>nnz</code> column indices, corresponding to the non-zero elements in the matrix

#### Output

<code>csrValA</code>	updated array of <code>nnz</code> elements, where <code>nnz</code> is the number of non-zero elements in the matrix
<code>csrRowPtrA</code>	updated array of <code>m+1</code> index elements
<code>csrColIndA</code>	updated array of <code>nnz</code> column indices, corresponding to the non-zero elements in the matrix

#### Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_ALLOC_FAILED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_INTERNAL_ERROR</b>	
<b>CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED</b>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}csr2dense

```

cusparseStatus_t
cusparseScsr2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const float *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    float *A, int lda )

cusparseStatus_t
cusparseDcsr2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const double *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    double *A, int lda )

cusparseStatus_t
cusparseCcsr2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const cuComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    cuComplex *A, int lda )

cusparseStatus_t
cusparseZcsr2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const cuDoubleComplex *csrValA,
    const int *csrRowPtrA, const int *csrColIndA,
    cuDoubleComplex *A, int lda )

```

Converts the matrix in CSR format defined by the three arrays `csrValA`, `csrRowPtrA`, and `csrColIndA` into a matrix `A` in dense format. The dense matrix `A` is filled in with the values of the sparse matrix and with zeros elsewhere.

### Input

<code>handle</code>	handle to a CUSPARSE context
<code>m</code>	number of rows of the matrix <code>A</code> ; <code>m</code> must be at least zero.
<code>n</code>	number of columns of matrix <code>A</code> ; <code>n</code> must be at least zero.

## Input (continued)

<code>descrA</code>	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
<code>cscValA</code>	array of nnz elements, where nnz is the number of non-zero elements and can be obtained from <code>cscRowPtrA(m) - cscRowPtrA(0)</code>
<code>cscRowPtrA</code>	array of m+1 index elements
<code>cscColIndA</code>	array of nnz column indices
<code>A</code>	array of dimension (lda, n)
<code>lda</code>	leading dimension of A

## Output

<code>A</code>	updated array filled in with values defined in the sparse matrix, and zeros elsewhere
----------------	---

Status Returned<sup>i</sup>

<b>CUSPARSE_STATUS_SUCCESS</b>	
<b>CUSPARSE_STATUS_NOT_INITIALIZED</b>	
<b>CUSPARSE_STATUS_INVALID_VALUE</b>	
<b>CUSPARSE_STATUS_ARCH_MISMATCH</b>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU
<b>CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED</b>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}dense2csc

```
cusparseStatus_t
cusparseSdense2csc(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const float *A, int lda,
    const int *nnzPerCol, float *cscValA,
    int *cscRowIndA, int *cscColPtrA )
```

```

cusparsesStatus_t
cusparsesDdense2csc(
    cusparsesHandle_t handle, int m, int n,
    const cusparsesMatDescr_t *descrA,
    const double *A, int lda,
    const int *nnzPerCol, double *cscValA,
    int *cscRowIndA, int *cscColPtrA )

cusparsesStatus_t
cusparsesCdense2csc(
    cusparsesHandle_t handle, int m, int n,
    const cusparsesMatDescr_t *descrA,
    const cuComplex *A, int lda,
    const int *nnzPerCol, cuComplex *cscValA,
    int *cscRowIndA, int *cscColPtrA )

cusparsesStatus_t
cusparsesZdense2csc(
    cusparsesHandle_t handle, int m, int n,
    const cusparsesMatDescr_t *descrA,
    const cuDoubleComplex *A, int lda,
    const int *nnzPerCol, cuDoubleComplex *cscValA,
    int *cscRowIndA, int *cscColPtrA )

```

Converts the matrix A in dense format into a matrix in CSC format. All the parameters are pre-allocated by the user, and the arrays are filled in based on nnzPerCol (which can be pre-computed with **cusparses{S,D,C,Z}nnz()**).

#### Input

handle	handle to a CUSPARSE context
m	number of rows of the matrix A; m must be at least zero.
n	number of columns of matrix A; n must be at least zero.
descrA	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
A	array of dimension (lda, n)
lda	leading dimension of A
nnzPerCol	array of size n containing the number of non-zero elements per column
cscValA	array of nnz elements to be filled

## Input (continued)

<code>cscRowIndA</code>	array of nnz row indices, corresponding to the non-zero elements in the matrix
<code>cscColPtrA</code>	array with n+1 index elements

## Output

<code>cscValA</code>	updated array of nnz elements, where nnz is the number of non-zero elements in the matrix
<code>cscRowIndA</code>	updated array of nnz row indices, corresponding to the non-zero elements in the matrix
<code>cscColPtrA</code>	updated array with n+1 index elements

Status Returned<sup>i</sup>

<code>CUSPARSE_STATUS_SUCCESS</code>	
<code>CUSPARSE_STATUS_NOT_INITIALIZED</code>	
<code>CUSPARSE_STATUS_ARCH_MISMATCH</code>	if the <b>D</b> or <b>Z</b> variants of the function were invoked on a device that does not support double precision.
<code>CUSPARSE_STATUS_EXECUTION_FAILED</code>	function failed to launch on GPU
<code>CUSPARSE_STATUS_MATRIX_TYPE_NOT_SUPPORTED</code>	

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}csc2dense

```

cusparseStatus_t
cusparseScsc2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const float *cscValA,
    const int *cscRowIndA, const int *cscColPtrA,
    float *A, int lda )

cusparseStatus_t
cusparseDcsc2dense(
    cusparseHandle_t handle, int m, int n,
    const cusparseMatDescr_t *descrA,
    const double *cscValA,
    const int *cscRowIndA, const int *cscColPtrA,
    double *A, int lda )

```

```

cusparsesStatus_t
cusparsesCcsc2dense(
    cusparsesHandle_t handle, int m, int n,
    const cusparsesMatDescr_t *descrA,
    const cuComplex *cscValA,
    const int *cscRowIndA, const int *cscColPtrA,
    cuComplex *A, int lda )

cusparsesStatus_t
cusparsesZcsc2dense(
    cusparsesHandle_t handle, int m, int n,
    const cusparsesMatDescr_t *descrA,
    const cuDoubleComplex *cscValA,
    const int *cscRowIndA, const int *cscColPtrA,
    cuDoubleComplex *A, int lda )

```

Converts the matrix in CSC format defined by the three arrays `cscValA`, `cscColPtrA`, and `cscRowIndA` into matrix `A` in dense format. The dense matrix `A` is filled in with the values of the sparse matrix and with zeros elsewhere.

#### Input

<code>handle</code>	handle to a CUSPARSE context
<code>m</code>	number of rows of the matrix <code>A</code> ; <code>m</code> must be at least zero.
<code>n</code>	number of columns of matrix <code>A</code> ; <code>n</code> must be at least zero.
<code>descrA</code>	descriptor of matrix <code>A</code> . The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
<code>cscValA</code>	array of <code>nnz</code> elements, where <code>nnz</code> is the number of non-zero elements and can be obtained from <code>cscColPtrA(m) - cscColPtrA(0)</code>
<code>cscRowIndA</code>	array of <code>nnz</code> row indices
<code>cscColPtrA</code>	array of <code>n+1</code> index elements
<code>A</code>	array of dimension <code>(lda, n)</code>
<code>lda</code>	leading dimension of <code>A</code>

#### Output

<code>A</code>	updated array filled in with values defined in the sparse matrix and zeros elsewhere
----------------	--

Status Returned<sup>i</sup>

---

CUSPARSE\_STATUS\_SUCCESS

CUSPARSE\_STATUS\_NOT\_INITIALIZED

CUSPARSE\_STATUS\_INVALID\_VALUE

CUSPARSE\_STATUS\_ARCH\_MISMATCH

if the D or Z variants of the function were invoked on a device that does not support double precision

CUSPARSE\_STATUS\_EXECUTION\_FAILED

function failed to launch on GPU

CUSPARSE\_STATUS\_MATRIX\_TYPE\_NOT\_SUPPORTED

---

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparse{S,D,C,Z}csr2csc

cusparseStatus\_t

cusparseScsr2csc(

    cusparseHandle\_t handle, int m, int n,  
    const float \*csrVal, const int \*csrRowPtr,  
    const int \*csrColInd, float \*cscVal,  
    int \*cscRowInd, int \*cscColPtr,  
    int copyValues, int base )

cusparseStatus\_t

cusparseDcsr2csc(

    cusparseHandle\_t handle, int m, int n,  
    const double \*csrVal,  
    const int \*csrRowPtr,  
    const int \*csrColInd, double \*cscVal,  
    int \*cscRowInd, int \*cscColPtr,  
    int copyValues, int base )

cusparseStatus\_t

cusparseCcsr2csc(

    cusparseHandle\_t handle, int m, int n,  
    const cuComplex \*csrVal, const int \*csrRowPtr,  
    const int \*csrColInd, cuComplex \*cscVal,  
    int \*cscRowInd, int \*cscColPtr,  
    int copyValues, int base )



```

cusparseStatus_t
cusparseZcsr2csc(
    cusparseHandle_t handle, int m, int n,
    const cuDoubleComplex *csrVal, const int *csrRowPtr,
    const int *csrColInd, cuDoubleComplex *cscVal,
    int *cscRowInd, int *cscColPtr,
    int copyValues, int base )

```

Converts the matrix in CSR format defined with the three arrays `csrVal`, `csrRowPtr`, and `csrColInd` into matrix A in CSC format defined by arrays `cscVal`, `cscRowInd`, and `cscColPtr`. The resulting matrix can also be seen as the transpose of the original sparse matrix. This routine can also be used to convert a matrix in CSC format into a matrix in CSR format.

#### Input

<code>handle</code>	handle to a CUSPARSE context
<code>m</code>	number of rows of the matrix A; <code>m</code> must be at least zero.
<code>n</code>	number of columns of matrix A; <code>n</code> must be at least zero.
<code>descrA</code>	descriptor of matrix A. The only <b>MatrixType</b> supported is <b>CUSPARSE_MATRIX_TYPE_GENERAL</b> . <b>IndexBase</b> constants <b>CUSPARSE_INDEX_BASE_ZERO</b> and <b>CUSPARSE_INDEX_BASE_ONE</b> are supported.
<code>csrVal</code>	array of <code>nnz</code> elements, where <code>nnz</code> is the number of non-zero elements and can be obtained from <code>csrRowPtrA(m) - csrRowPtrA(0)</code>
<code>csrRowPtr</code>	array of <code>m+1</code> indices
<code>csrColInd</code>	array of <code>nnz</code> column indices
<code>cscVal</code>	array of <code>nnz</code> elements, where <code>nnz</code> is the number of non-zero elements and can be obtained from <code>csrColPtr(m) - csrColPtr(0)</code>
<code>cscRowInd</code>	array of <code>nnz</code> row indices
<code>cscColPtr</code>	array of <code>n+1</code> indices
<code>copyValues</code>	if zero, <code>cscVal</code> array is not filled
<code>base</code>	base index: <b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

#### Output

<code>cscVal</code>	if <code>copyValues</code> is non-zero, updated array
<code>cscColPtr</code>	updated array of <code>n+1</code> index elements
<code>cscRowInd</code>	updated array of <code>nnz</code> row indices,

Status Returned<sup>i</sup>

---

**CUSPARSE\_STATUS\_SUCCESS**

**CUSPARSE\_STATUS\_NOT\_INITIALIZED**

**CUSPARSE\_STATUS\_ALLOC\_FAILED**

**CUSPARSE\_STATUS\_INVALID\_VALUE**

**CUSPARSE\_STATUS\_ARCH\_MISMATCH**

if the **D** or **Z** variants of the function were invoked on a device that does not support double precision.

**CUSPARSE\_STATUS\_EXECUTION\_FAILED**

function failed to launch on GPU

**CUSPARSE\_STATUS\_INTERNAL\_ERROR**

---

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseXcoo2csr

**cusparseStatus\_t**

**cusparseXcoo2csr(**

**cusparseHandle\_t handle, const int \*cooRowInd,**  
**int nnz, int m, int \*csrRowPtr,**  
**cusparseIndexBase\_t idxBase )**

Converts the array containing the uncompressed row indices (corresponding to COO format) into an array of compressed row pointers (corresponding to CSR format).

It can also be used to convert the array containing the uncompressed column indices (corresponding to COO format) into an array of column pointers (corresponding to CSC format).

Input

---

handle	handle to a CUSPARSE context
cooRowInd	array of row indices
nnz	number of non-zeros of the matrix in COO format; this is also the length of array cooRowInd
m	number of rows of the matrix A; m must be at least zero.
csrRowPtr	array of row pointers
idxBase	base index: <b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

---

Output

---

csrRowPtr	updated array of m+1 index elements
-----------	-------------------------------------

---

Status Returned<sup>i</sup>

---

**CUSPARSE\_STATUS\_SUCCESS**

**CUSPARSE\_STATUS\_NOT\_INITIALIZED**

**CUSPARSE\_STATUS\_INVALID\_VALUE**

if idxBase is neither

**CUSPARSE\_INDEX\_BASE\_ZERO** nor

**CUSPARSE\_INDEX\_BASE\_ONE**

**CUSPARSE\_STATUS\_EXECUTION\_FAILED**

function failed to launch on GPU

---

i. See also [CUSPARSE Status Definitions](#) on page 14.

## cusparseXcsr2coo

**cusparseStatus\_t**

**cusparseXcsr2coo(**

**cusparseHandle\_t handle, const int \*csrRowPtr,**

**int nnz, int m, int \*cooRowInd,**

**cusparseIndexBase\_t idxBase )**

Converts the array containing the compressed row pointers (corresponding to CSR format) into an array of uncompressed row indices (corresponding to COO format).

It can also be used to convert the array containing the compressed column pointers (corresponding to CSC format) into an array of uncompressed column indices (corresponding to COO format).

Input

---

handle	handle to a CUSPARSE context
csrRowPtr	array of compressed row pointers
nnz	number of non-zeros of the matrix in COO format; this is also the length of array cooRowInd
m	number of rows of the matrix A; m must be at least zero.
cooRowInd	array of uncompressed row indices
idxBase	base index: <b>CUSPARSE_INDEX_BASE_ZERO</b> or <b>CUSPARSE_INDEX_BASE_ONE</b>

---

Output

---

cooRowInd	updated array of nnz index elements
-----------	-------------------------------------

---

Status Returned<sup>i</sup>

---

**CUSPARSE\_STATUS\_SUCCESS**

**CUSPARSE\_STATUS\_NOT\_INITIALIZED**

Status Returned<sup>i</sup> (continued)

<b>CUSPARSE_STATUS_INVALID_VALUE</b>	if <code>idxBase</code> is neither <b>CUSPARSE_INDEX_BASE_ZERO</b> nor <b>CUSPARSE_INDEX_BASE_ONE</b>
<b>CUSPARSE_STATUS_EXECUTION_FAILED</b>	function failed to launch on GPU

i. See also [CUSPARSE Status Definitions](#) on page 14.

# A

## CUSPARSE Library Example

[Example 1](#) on page 54 demonstrates an application of the CUSPARSE library. The example performs these actions:

1. Creates a sparse test matrix in COO format.
2. Creates a sparse and dense vector.
3. Allocates GPU memory and copies the matrix and vectors into it.
4. Initializes the CUSPARSE library.
5. Creates and sets up the matrix descriptor.
6. Converts the matrix from COO to CSR format.
7. Exercises Level 1 routines.
8. Exercises Level 2 routines.
9. Exercises Level 3 routines.

### Example 1. Using the CUSPARSE Library

---

```

#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cusparse.h"

#define CLEANUP(s) \
do { \
    printf ("%s\n", s); \
    if (yHostPtr) free(yHostPtr); \
    if (zHostPtr) free(zHostPtr); \
    if (xIndHostPtr) free(xIndHostPtr); \
    if (xValHostPtr) free(xValHostPtr); \
    if (cooRowIndexHostPtr) free(cooRowIndexHostPtr); \
    if (cooColIndexHostPtr) free(cooColIndexHostPtr); \
    if (cooValHostPtr) free(cooValHostPtr); \
    if (y) cudaFree(y); \
    if (z) cudaFree(z); \
    if (xInd) cudaFree(xInd); \
    if (xVal) cudaFree(xVal); \
    if (csrRowPtr) cudaFree(csrRowPtr); \
    if (cooRowIndex) cudaFree(cooRowIndex); \
    if (cooColIndex) cudaFree(cooColIndex); \
    if (cooVal) cudaFree(cooVal); \
    if (handle) cusparseDestroy(handle); \
    fflush (stdout); \
} while (0)

int main(){
    cudaError_t cudaStat1,cudaStat2,cudaStat3,cudaStat4,cudaStat5,cudaStat6;
    cusparseStatus_t status;
    cusparseHandle_t handle=0;
    cusparseMatDescr_t descra=0;
    int * cooRowIndexHostPtr=0;
    int * cooColIndexHostPtr=0;
    double * cooValHostPtr=0;
    int * cooRowIndex=0;

```

## Example 1. Using the CUSPARSE Library (continued)

---

```

int *    cooColIndex=0;
double * cooVal=0;
int *    xIndHostPtr=0;
double * xValHostPtr=0;
double * yHostPtr=0;
int *    xInd=0;
double * xVal=0;
double * y=0;
int *    csrRowPtr=0;
double * zHostPtr=0;
double * z=0;
int      n, nnz, nnz_vector, i, j;

printf("testing example\n");
/* create the following sparse test matrix in COO format */
/* |1.0      2.0 3.0|
   |  4.0      |
   |5.0      6.0 7.0|
   |  8.0      9.0| */
n=4; nnz=9;
cooRowIndexHostPtr = (int *) malloc(nnz*sizeof(cooRowIndexHostPtr[0]));
cooColIndexHostPtr = (int *) malloc(nnz*sizeof(cooColIndexHostPtr[0]));
cooValHostPtr      = (double *)malloc(nnz*sizeof(cooValHostPtr[0]));
if ((!cooRowIndexHostPtr) || (!cooColIndexHostPtr) || (!cooValHostPtr)){
    CLEANUP("Host malloc failed (matrix)");
    return EXIT_FAILURE;
}
cooRowIndexHostPtr[0]=0; cooColIndexHostPtr[0]=0; cooValHostPtr[0]=1.0;
cooRowIndexHostPtr[1]=0; cooColIndexHostPtr[1]=2; cooValHostPtr[1]=2.0;
cooRowIndexHostPtr[2]=0; cooColIndexHostPtr[2]=3; cooValHostPtr[2]=3.0;
cooRowIndexHostPtr[3]=1; cooColIndexHostPtr[3]=1; cooValHostPtr[3]=4.0;
cooRowIndexHostPtr[4]=2; cooColIndexHostPtr[4]=0; cooValHostPtr[4]=5.0;
cooRowIndexHostPtr[5]=2; cooColIndexHostPtr[5]=2; cooValHostPtr[5]=6.0;
cooRowIndexHostPtr[6]=2; cooColIndexHostPtr[6]=3; cooValHostPtr[6]=7.0;
cooRowIndexHostPtr[7]=3; cooColIndexHostPtr[7]=1; cooValHostPtr[7]=8.0;
cooRowIndexHostPtr[8]=3; cooColIndexHostPtr[8]=3; cooValHostPtr[8]=9.0;

```

### Example 1. Using the CUSPARSE Library (continued)

---

```
//print the matrix
printf("Input data:\n");
for (i=0; i<nnz; i++){
    printf("cooRowIndexHostPtr[%d]=%d  ",i,cooRowIndexHostPtr[i]);
    printf("cooColIndexHostPtr[%d]=%d  ",i,cooColIndexHostPtr[i]);
    printf("cooValHostPtr[%d]=%f      \n",i,cooValHostPtr[i]);
}

/* create a sparse and dense vector */
/* xVal= [100.0 200.0 400.0]    (sparse)
   xInd= [0      1      3      ]
   y    = [10.0 20.0 30.0 40.0 | 50.0 60.0 70.0 80.0] (dense) */
nnz_vector = 3;
xIndHostPtr = (int *)    malloc(nnz_vector*sizeof(xIndHostPtr[0]));
xValHostPtr = (double *)malloc(nnz_vector*sizeof(xValHostPtr[0]));
yHostPtr    = (double *)malloc(2*n      *sizeof(yHostPtr[0]));
zHostPtr    = (double *)malloc(2*(n+1)  *sizeof(zHostPtr[0]));
if((!xIndHostPtr) || (!xValHostPtr) || (!yHostPtr) || (!zHostPtr)){
    CLEANUP("Host malloc failed (vectors)");
    return EXIT_FAILURE;
}

yHostPtr[0] = 10.0;  xIndHostPtr[0]=0; xValHostPtr[0]=100.0;
yHostPtr[1] = 20.0;  xIndHostPtr[1]=1; xValHostPtr[1]=200.0;
yHostPtr[2] = 30.0;
yHostPtr[3] = 40.0;  xIndHostPtr[2]=3; xValHostPtr[2]=400.0;
yHostPtr[4] = 50.0;
yHostPtr[5] = 60.0;
yHostPtr[6] = 70.0;
yHostPtr[7] = 80.0;
//print the vectors
for (j=0; j<2; j++){
    for (i=0; i<n; i++){
        printf("yHostPtr[%d,%d]=%f\n",i,j,yHostPtr[i+n*j]);
    }
}

for (i=0; i<nnz_vector; i++){
    printf("xIndHostPtr[%d]=%d  ",i,xIndHostPtr[i]);
}
```



## Example 1. Using the CUSPARSE Library (continued)

---

```

    printf("xValHostPtr[%d]=%f\n",i,xValHostPtr[i]);
}

/* allocate GPU memory and copy the matrix and vectors into it */
cudaStat1 = cudaMalloc((void**)&cooRowIndex,nnz*sizeof(cooRowIndex[0]));
cudaStat2 = cudaMalloc((void**)&cooColIndex,nnz*sizeof(cooColIndex[0]));
cudaStat3 = cudaMalloc((void**)&cooVal,      nnz*sizeof(cooVal[0]));
cudaStat4 = cudaMalloc((void**)&y,          2*n*sizeof(y[0]));
cudaStat5 = cudaMalloc((void**)&xInd,nnz_vector*sizeof(xInd[0]));
cudaStat6 = cudaMalloc((void**)&xVal,nnz_vector*sizeof(xVal[0]));
if ((cudaStat1 != cudaSuccess) ||
    (cudaStat2 != cudaSuccess) ||
    (cudaStat3 != cudaSuccess) ||
    (cudaStat4 != cudaSuccess) ||
    (cudaStat5 != cudaSuccess) ||
    (cudaStat6 != cudaSuccess)) {
    CLEANUP("Device malloc failed");
    return EXIT_FAILURE;
}
cudaStat1 = cudaMemcpy(cooRowIndex, cooRowIndexHostPtr,
                      (size_t)(nnz*sizeof(cooRowIndex[0])),
                      cudaMemcpyHostToDevice);
cudaStat2 = cudaMemcpy(cooColIndex, cooColIndexHostPtr,
                      (size_t)(nnz*sizeof(cooColIndex[0])),
                      cudaMemcpyHostToDevice);
cudaStat3 = cudaMemcpy(cooVal,      cooValHostPtr,
                      (size_t)(nnz*sizeof(cooVal[0])),
                      cudaMemcpyHostToDevice);
cudaStat4 = cudaMemcpy(y,          yHostPtr,
                      (size_t)(2*n*sizeof(y[0])),
                      cudaMemcpyHostToDevice);
cudaStat5 = cudaMemcpy(xInd,      xIndHostPtr,
                      (size_t)(nnz_vector*sizeof(xInd[0])),
                      cudaMemcpyHostToDevice);
cudaStat6 = cudaMemcpy(xVal,      xValHostPtr,
                      (size_t)(nnz_vector*sizeof(xVal[0])),
                      cudaMemcpyHostToDevice);

```

## Example 1. Using the CUSPARSE Library (continued)

---

```

    if ((cudaStat1 != cudaSuccess) ||
        (cudaStat2 != cudaSuccess) ||
        (cudaStat3 != cudaSuccess) ||
        (cudaStat4 != cudaSuccess) ||
        (cudaStat5 != cudaSuccess) ||
        (cudaStat6 != cudaSuccess)) {
        CLEANUP("Memcpy from Host to Device failed");
        return EXIT_FAILURE;
    }

    /* initialize cusparse library */
    status= cusparseCreate(&handle);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("CUSPARSE Library initialization failed");
        return EXIT_FAILURE;
    }

    /* create and setup matrix descriptor */
    status= cusparseCreateMatDescr(&descra);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("Matrix descriptor initialization failed");
        return EXIT_FAILURE;
    }
    cusparseSetMatType(descra,CUSPARSE_MATRIX_TYPE_GENERAL);
    cusparseSetMatIndexBase(descra,CUSPARSE_INDEX_BASE_ZERO);

    /* exercise conversion routines (convert matrix from COO 2 CSR format) */
    cudaStat1 = cudaMalloc((void**)&csrRowPtr,(n+1)*sizeof(csrRowPtr[0]));
    if (cudaStat1 != cudaSuccess) {
        CLEANUP("Device malloc failed (csrRowPtr)");
        return EXIT_FAILURE;
    }
    status= cusparseXcoo2csr(handle,cooRowIndex,nnz,n,
                            csrRowPtr,CUSPARSE_INDEX_BASE_ZERO);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("Conversion from COO to CSR format failed");
    }

```

## Example 1. Using the CUSPARSE Library (continued)

---

```

        return EXIT_FAILURE;
    }
    //csrRowPtr = [0 3 4 7 9]

    /* exercise Level 1 routines (scatter vector elements) */
    status= cusparseDscctr(handle, nnz_vector, xVal, xInd,
                          &y[n], CUSPARSE_INDEX_BASE_ZERO);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("Scatter from sparse to dense vector failed");
        return EXIT_FAILURE;
    }
    //y = [10 20 30 40 | 100 200 70 400]

    /* exercise Level 2 routines (csrcmv) */
    status= cusparseDcsrcmv(handle,CUSPARSE_OPERATION_NON_TRANSPOSE, n, n, 2.0,
                          descra, cooVal, csrRowPtr, cooColIndex, &y[0],
                          3.0, &y[n]);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("Matrix-vector multiplication failed");
        return EXIT_FAILURE;
    }

    /* print intermediate results (y) */
    //y = [10 20 30 40 | 680 760 1230 2240]
    cudaMemcpy(yHostPtr, y, (size_t)(2*n*sizeof(y[0])), cudaMemcpyDeviceToHost);
    printf("Intermediate results:\n");
    for (j=0; j<2; j++){
        for (i=0; i<n; i++){
            printf("yHostPtr[%d,%d]=%f\n",i,j,yHostPtr[i+n*j]);
        }
    }

    /* exercise Level 3 routines (csrmm) */
    cudaStat1 = cudaMalloc((void**)&z, 2*(n+1)*sizeof(z[0]));
    if (cudaStat1 != cudaSuccess) {
        CLEANUP("Device malloc failed (z)");
        return EXIT_FAILURE;
    }

```

### Example 1. Using the CUSPARSE Library (continued)

---

```

    }
    cudaStat1 = cudaMemset((void *)z,0, 2*(n+1)*sizeof(z[0]));
    if (cudaStat1 != cudaSuccess) {
        CLEANUP("Memset on Device failed");
        return EXIT_FAILURE;
    }
    status= cusparseDcsrmm(handle, CUSPARSE_OPERATION_NON_TRANSPOSE, n, 2, n,
                          5.0, descra, cooVal, csrRowPtr, cooColIndex, y, n,
                          0.0, z, n+1);
    if (status != CUSPARSE_STATUS_SUCCESS) {
        CLEANUP("Matrix-matrix multiplication failed");
        return EXIT_FAILURE;
    }

    /* print final results (z) */
    cudaStat1 = cudaMemcpy(zHostPtr, z,
                          (size_t)(2*(n+1)*sizeof(z[0])),
                          cudaMemcpyDeviceToHost);
    if (cudaStat1 != cudaSuccess) {
        CLEANUP("Memcpy from Device to Host failed");
        return EXIT_FAILURE;
    }
    //z = [950 400 2550 2600 0 | 49300 15200 132300 131200 0]
    printf("Final results:\n");
    for (j=0; j<2; j++){
        for (i=0; i<n+1; i++){
            printf("z[%d,%d]=%f\n",i,j,zHostPtr[i+(n+1)*j]);
        }
    }

    /* check the results */
    /* Notice that CLEANUP() contains a call to cusparseDestroy(handle) */
    if ((zHostPtr[0] != 950.0)    ||
        (zHostPtr[1] != 400.0)    ||
        (zHostPtr[2] != 2550.0)   ||
        (zHostPtr[3] != 2600.0)   ||
        (zHostPtr[4] != 0.0)     ||

```

## Example 1. Using the CUSPARSE Library (continued)

---

```
(zHostPtr[5] != 49300.0) ||
(zHostPtr[6] != 15200.0) ||
(zHostPtr[7] != 132300.0) ||
(zHostPtr[8] != 131200.0) ||
(zHostPtr[9] != 0.0)      ||
(yHostPtr[0] != 10.0)     ||
(yHostPtr[1] != 20.0)     ||
(yHostPtr[2] != 30.0)     ||
(yHostPtr[3] != 40.0)     ||
(yHostPtr[4] != 680.0)    ||
(yHostPtr[5] != 760.0)    ||
(yHostPtr[6] != 1230.0)   ||
(yHostPtr[7] != 2240.0)){
    CLEANUP("example test FAILED");
    return EXIT_FAILURE;
}
else{
    CLEANUP("example test PASSED");
    return EXIT_SUCCESS;
}
}
```

---